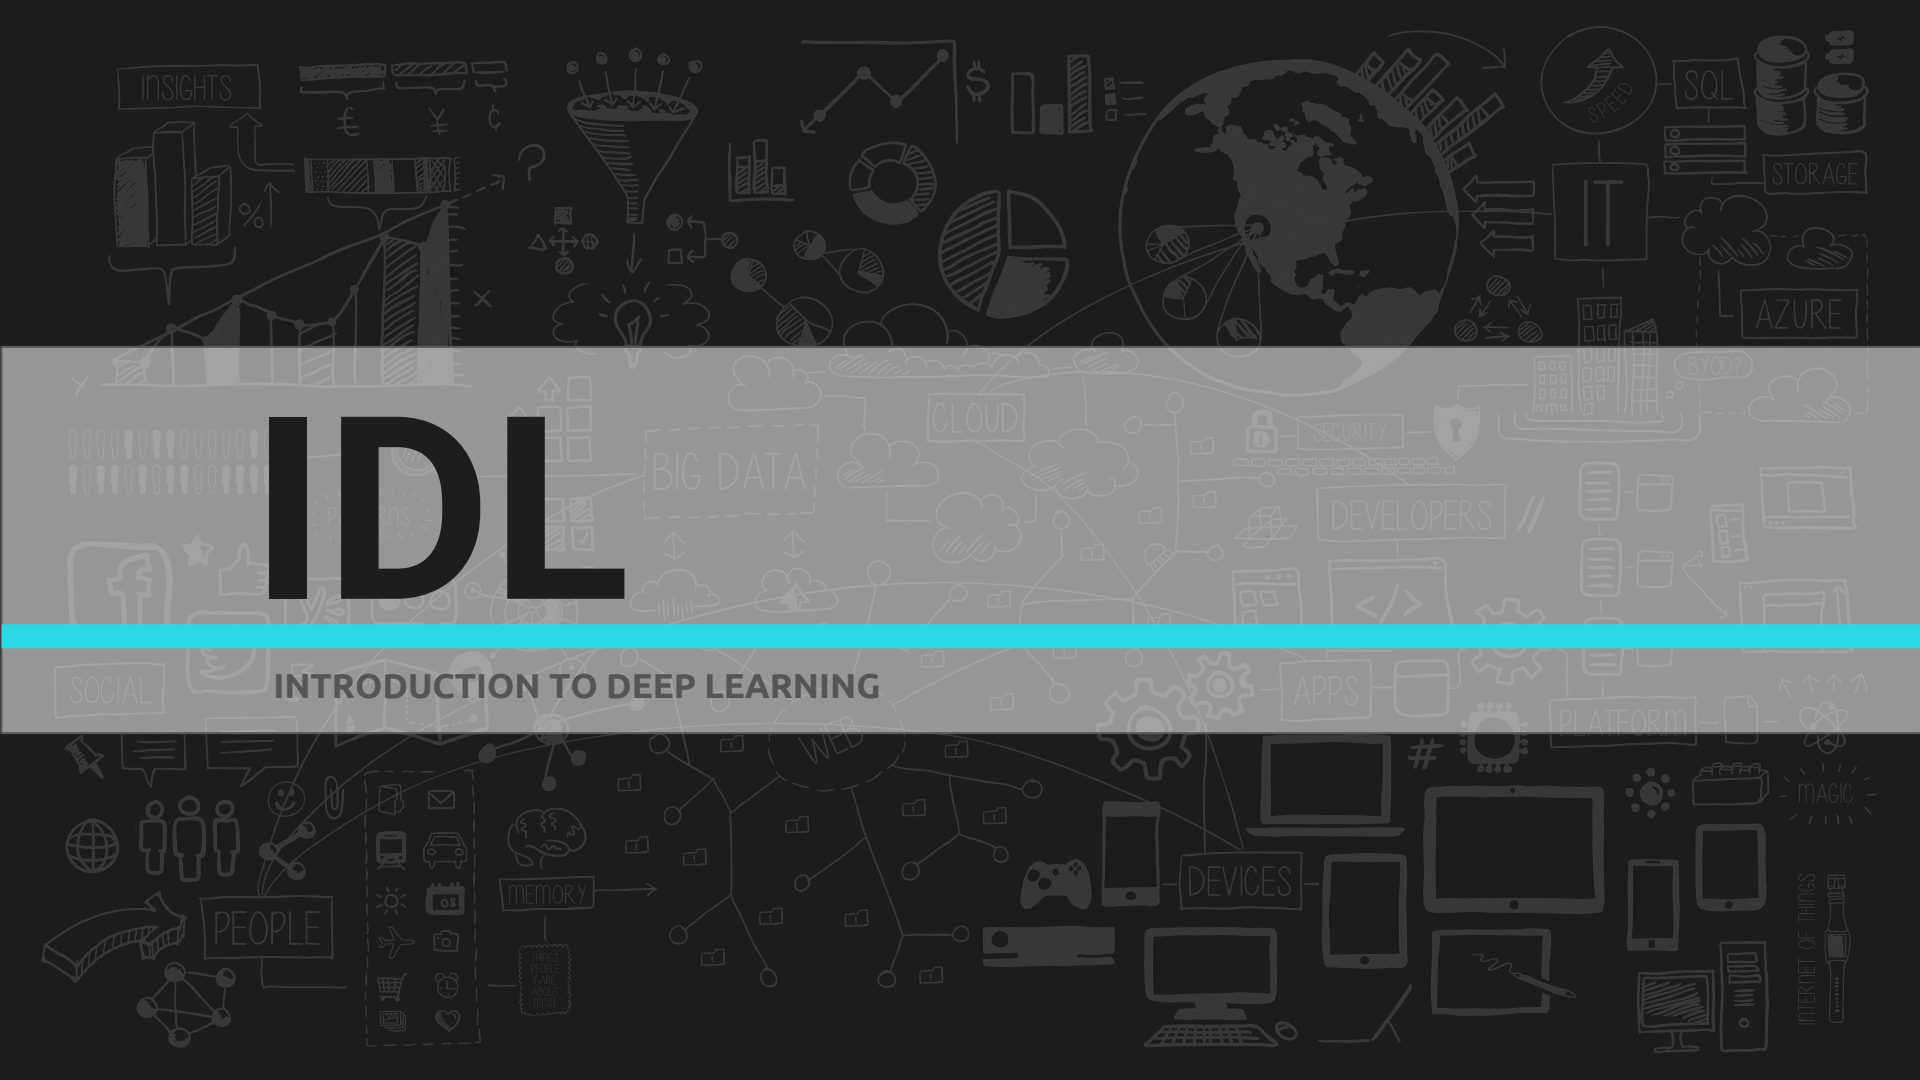

IDL

INTRODUCTION TO DEEP LEARNING

Optimization

- Mini-batch gradient descent

- Optimizer **:**
    - SGD
    - GD with momentum
    - RMSProp
    - AdaGrad
    - Adam

- Learning rate decay

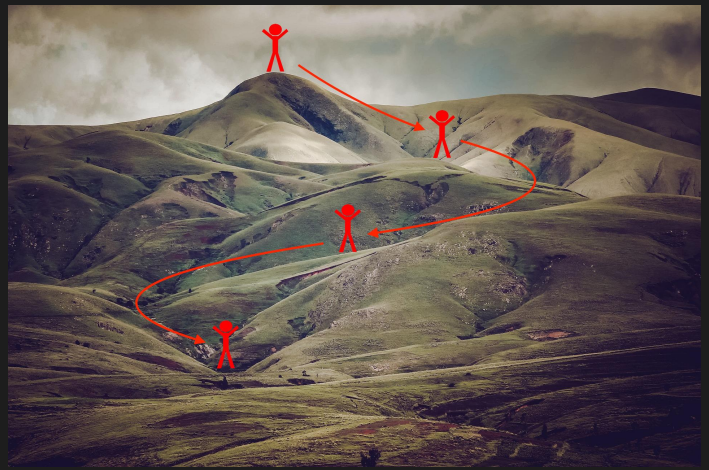

Mini-batch gradient descent

Exponentially weighted averages

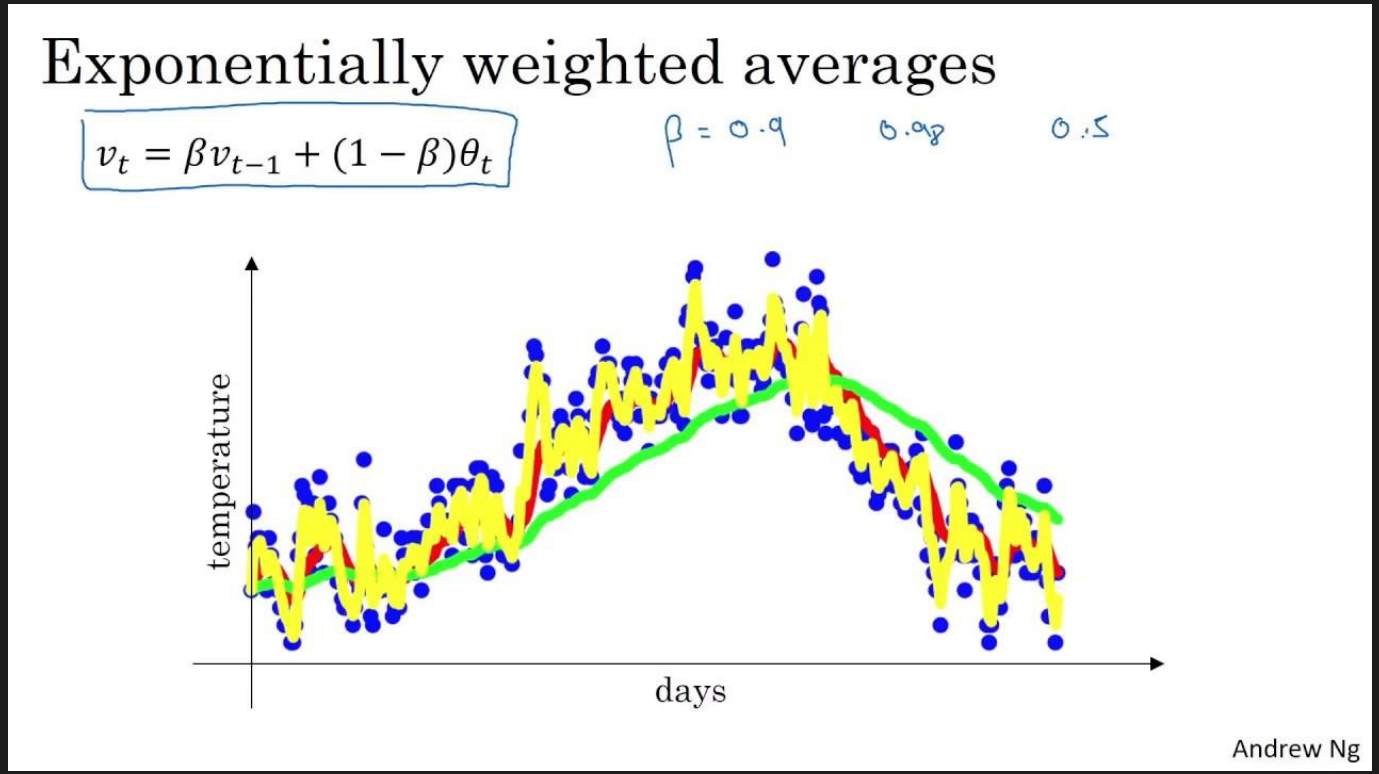

GD with momentum

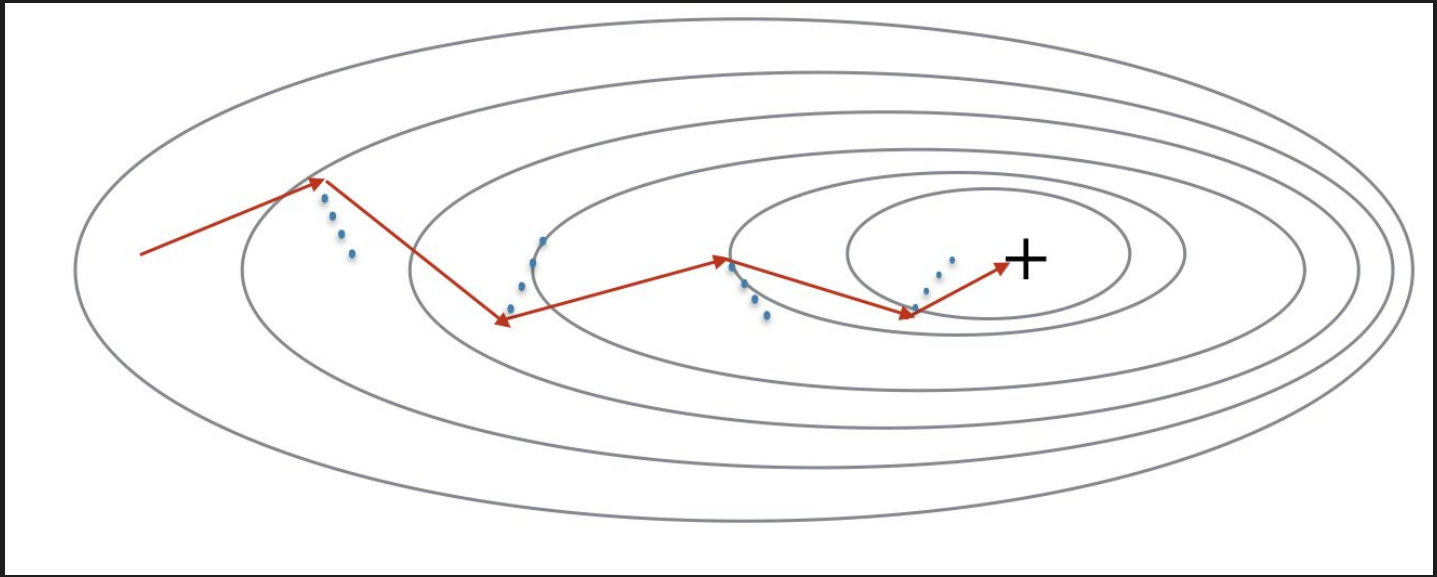

RMSProp

$$v_{dw} = \beta \cdot v_{dw} + (1 - \beta) \cdot dw^2$$

$$v_{db} = \beta \cdot v_{dw} + (1 - \beta) \cdot db^2$$

$$W = W - \alpha \cdot \frac{dw}{\sqrt{v_{dw}} + \epsilon}$$

$$b = b - \alpha \cdot \frac{db}{\sqrt{v_{db}} + \epsilon}$$

Adam

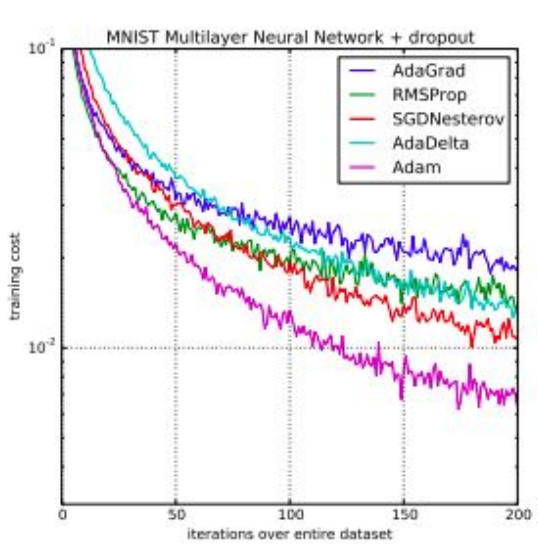

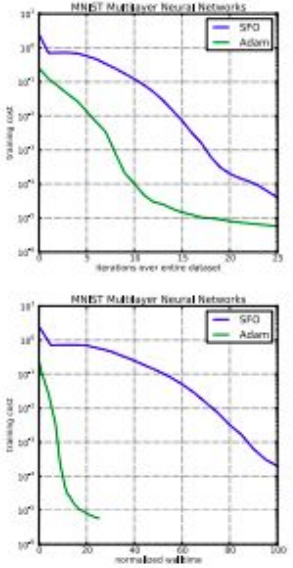

(a) (b)

Figure 2: Training of multilayer neural networks on MNIST images. (a) Neural networks using dropout stochastic regularization. (b) Neural networks with deterministic cost function. We compare with the sum-of-functions (SFO) optimizer (Sohl-Dickstein et al., 2014)

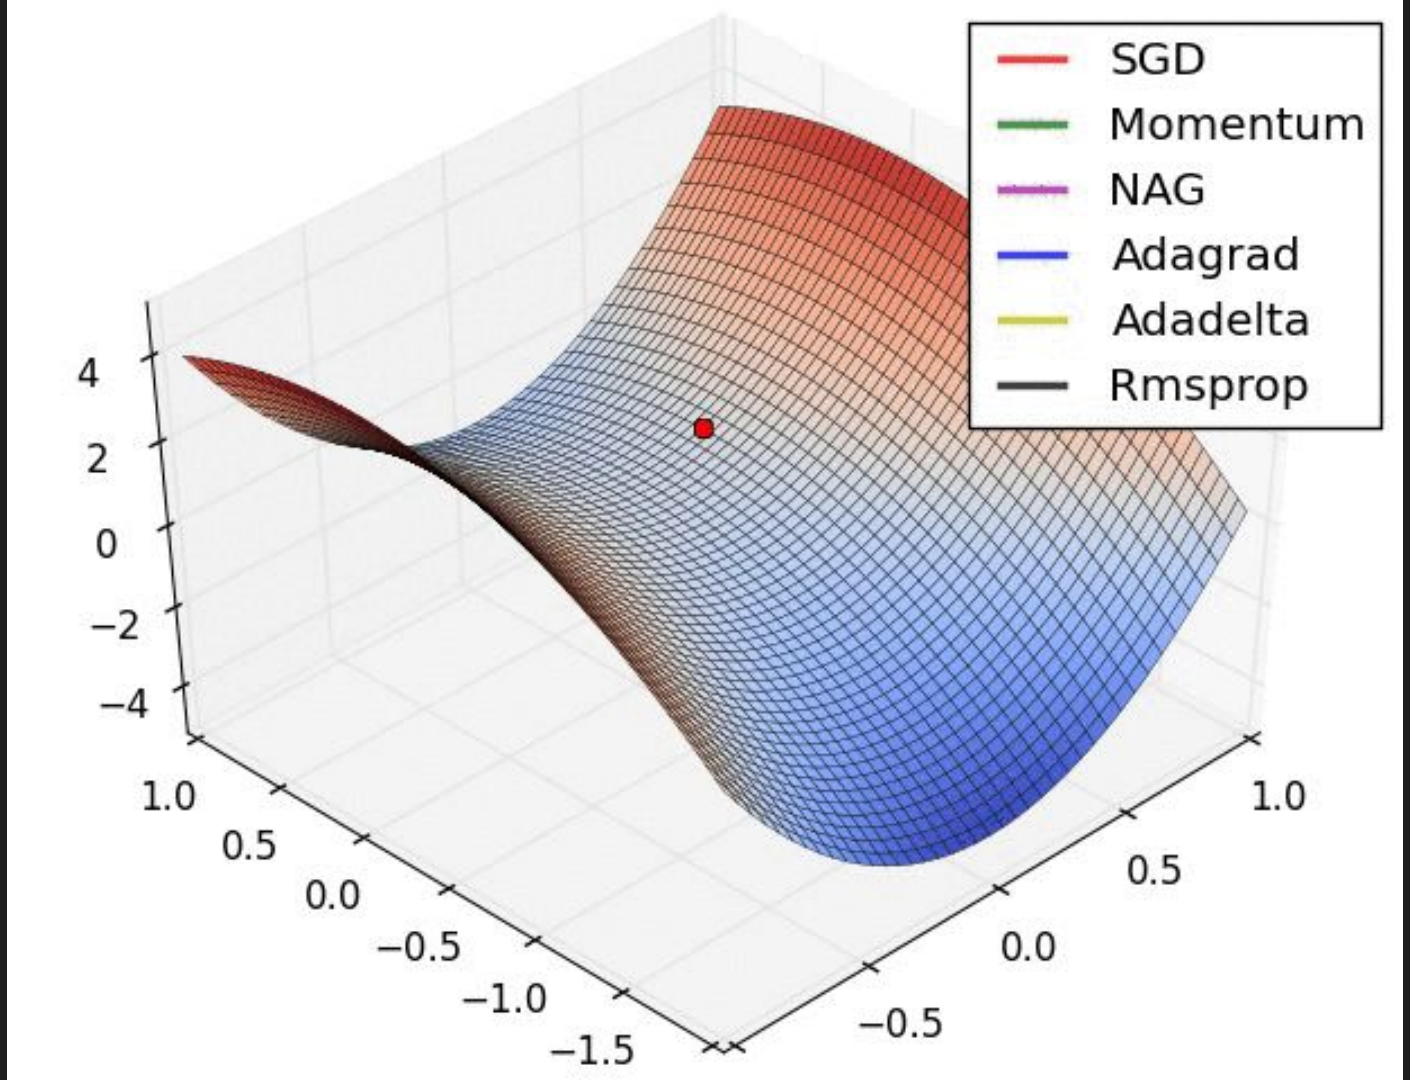
SGD
Momentum
NAG
Adagrad
Adadelta
Rmsprop
4
2
0
−2
−4
1.0
0.5
0.0
−0.5
−1.0
−1.5
1.0
0.5
0.0
−0.5

The problem of local optima

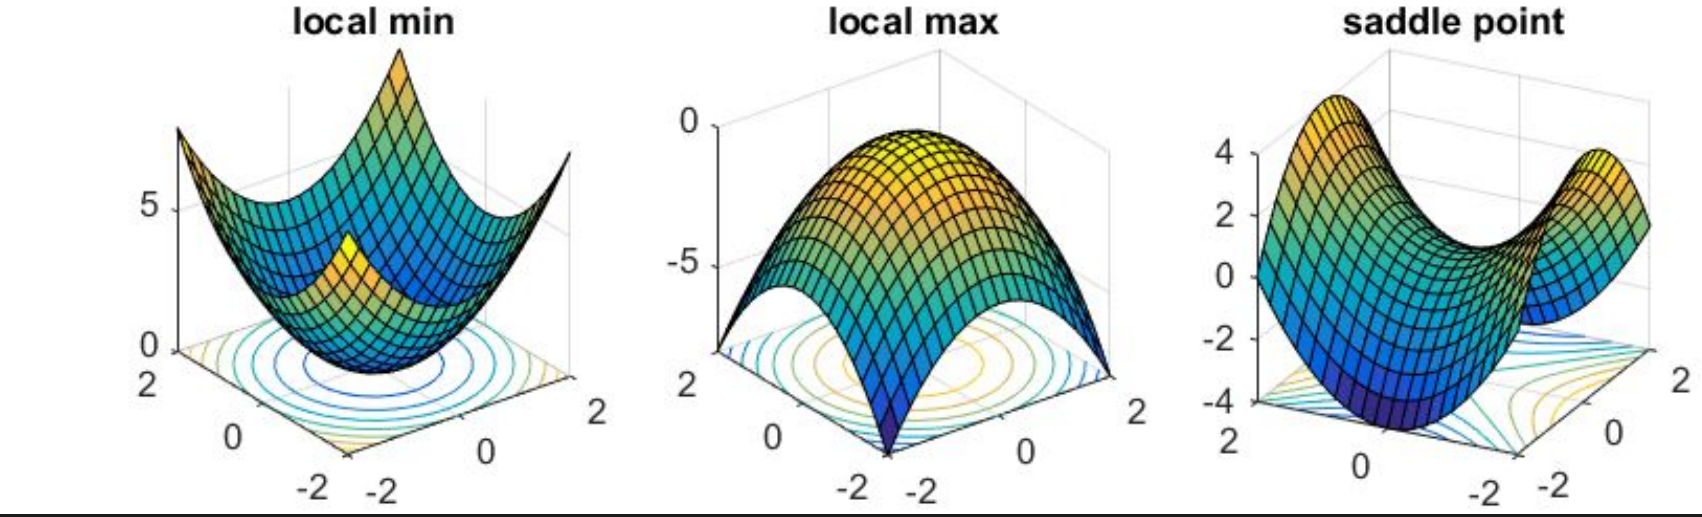

Regularization

Hyperparameter

Batch Normalization